

Transaction concept:

Unit of program execution that access and updates various data items.

e.g.: transaction of \$50 from account A → B

1) read(A)

2) A := A - 50

3) write(A)

4) read(B)

5) B := B + 50

6) write(B)

Two main issues to deal with:

→ Failure of various kinds, such as hardware failure, system crash

→ concurrent execution: multiple "transact" in same account at same instant

#

Required properties of ACID

1) Atomicity: means 0 or 1

either complete "transaction" or no transaction

$\{ \text{ACI} \}$ if for A, it did not reflect then it's not reflect

2) Consistency:

→ A + B must unchanged by execution of transaction

→ sum of amount before transaction equal to sum of transaction after execution.

being alone
3) ISOLATION: used when two/more transactⁿ happen in some o/p at same time

→ if b/w step 3 + 6 of transactⁿ, another transactⁿ T₂ happens is allowed to access the partially updated database it will see inconsistent.

→ every Transactⁿ are independent.

→ Isolatⁿ can be ensured trivially by running transactⁿ serially.

4) DURABILITY:

→ once user get notified that transactⁿ completed, the update to the database by the transaction must persist even if System Crash.

→ Someone send \$10

→ my balance now \$60 (earlier \$50)
System Crash

→ after system updated, my a/c balance must show \$60.

A → ATOMICITY

(All or Nothing Transactⁿ)

C - consistency

(A+B remain constant)

I - Isolation

(Transactⁿs are independent)

D - Durability

(committed data never lost)

TRANSACTION STATE

→ Active : initial state

→ Partially committed

→ Failed

→ Aborted

↳ after transaction has been rolled back (refund)

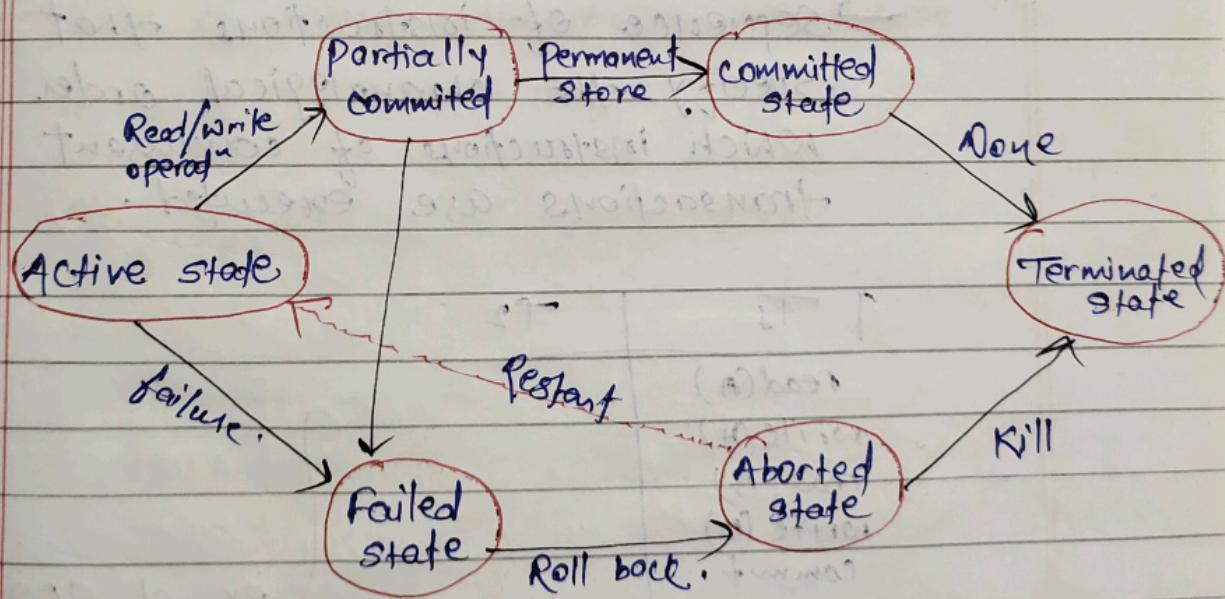
Two option after aborted :-

i) Restart the transaction

ii) Kill the transaction.

→ Committed : after successful completion

→ Terminated : after committed.



Concurrent Executions

→ Multiple transactions are allowed to run concurrently in systems. Advantages are :-

- i) Increase processor & disc utilization
- ii) Reduce av. response time.

(bwpri) Concurrency control mechanism to achieve isolation.

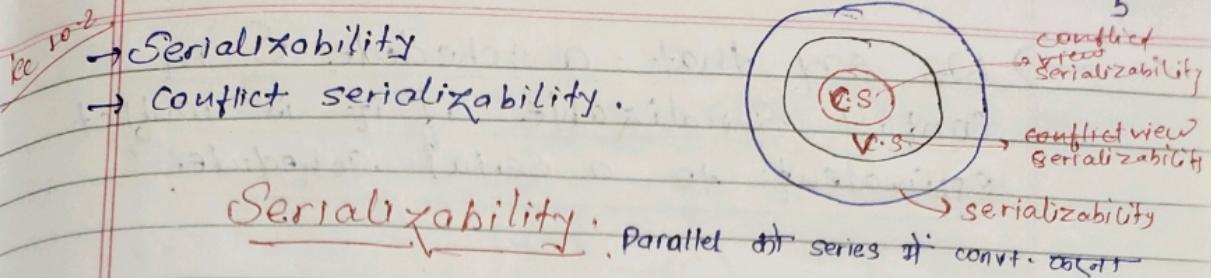
Schedule: two or more transaction together called schedule.

→ Sequence of instructions that specify the chronological order in which instructions of concurrent transactions are executed.

T ₁	T ₂
read(A)	
write(A)	
read(B)	
write(B)	
commit	
	read(A)
	write(A)
	read(B)
	write(B)
	commit

} ex. of SERIAL Schedule

if in serial schedule if there are 'n' transaction then total $n!$ equivalent serial schedule possible



Assumpt: each transaction preserve database consistency

- \rightarrow Conflict serializability } parallel or series
- \rightarrow View serializability } if convt. or not or 2 method

CONFLICT Instruction

- \rightarrow Two transaction T_1 & T_2
- $\rightarrow T_1$ & T_2 conflict if and only if (at least one of these instructions is write)

T_1	T_2	T_1	T_2	T_1	T_2	T_1	T_2
R(A)	.	R(A)	.	W(A)	.	R(A)	W(A)
R(B)	.	W(B)	.	W(A)	.	W(B)	W(B)

Not conflict conflict conflict conflict

Conflict Serializability only if graph is acyclic

- \hookrightarrow concept that ensure transactions can be executed in a way that maintains the same result as if they were executed one after another.

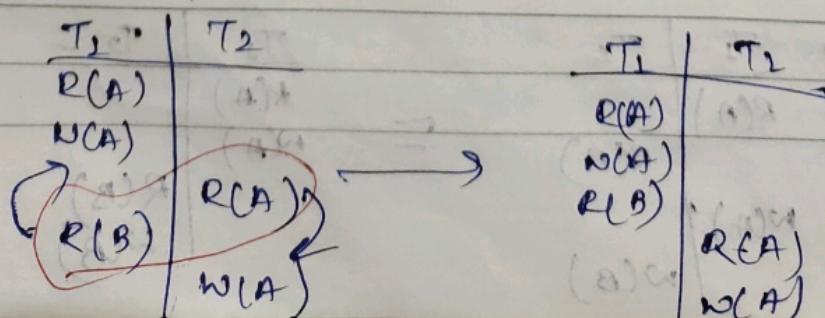
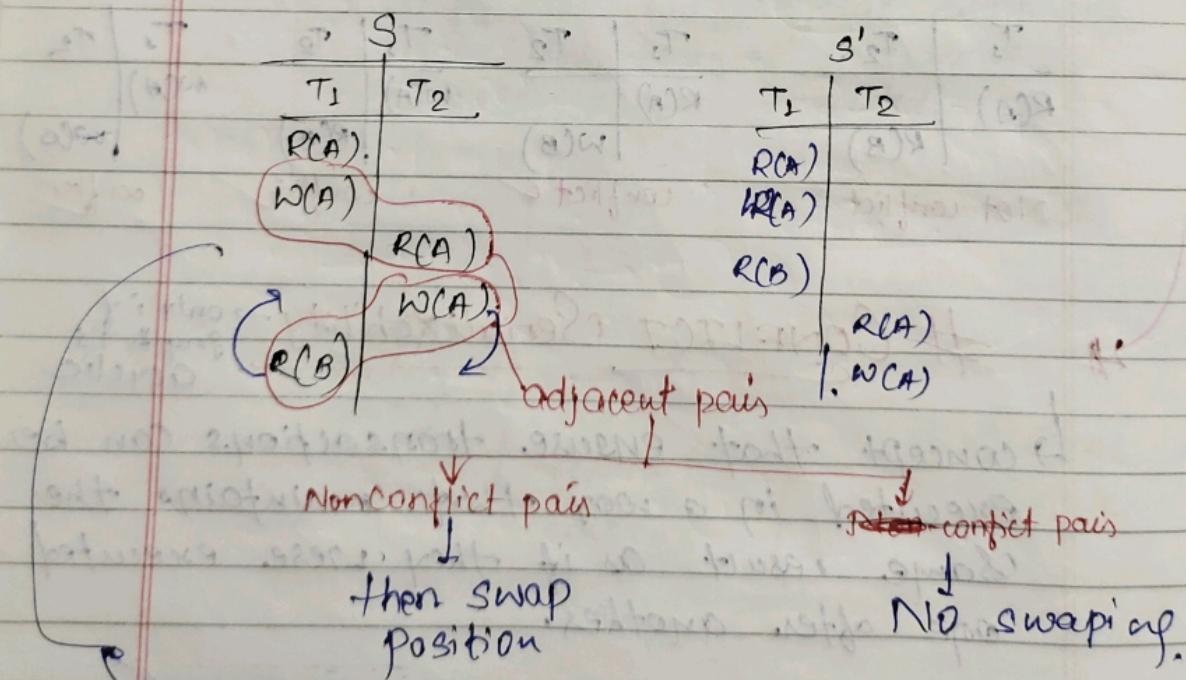
T_1	T_2	T_3	T_2
R(A)	.	R(A)	.
R(B)	.	W(A)	.
W(A)	.	R(B)	W(B)
W(B)	.	W(B)	W(B)

→ We say that a schedule S is conflict serializable if it is conflict equivalent to a serial schedule.

$R(A), R(A)$ } non-conflict pairs
 R

$R(A), W(A)$ }
 $W(A), R(A)$ } conflict pairs
 $W(A), W(A)$

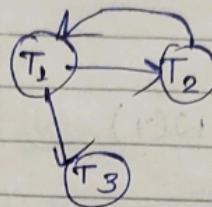
$R(B), R(A)$ }
 $W(B), R(A)$ }
 $R(B), W(A)$ } Non-conflict pairs
 $W(A), W(B)$ }
 $\xrightarrow{\text{string of non-conflict pairs}}$



in conflict
i.e.

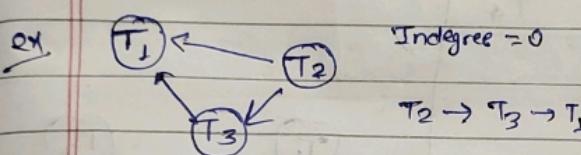
S		
T ₁	T ₂	T ₃
R(A)	W(A)	
W(A)		W(A)

Check whether schedule is conflict serializable or not?



अगर graph में
loop आ जाए
means Non-conflict
serializable.

जिस तिले pair हैं
conflict pair बना
जाता है उसे → से
join कर दो



No loop/cycle
↓
Conflict serializable
↓
serial
consistent

परंतु loop आ जाता है तो → Serializability

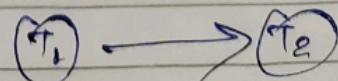
Check next in लिए View serializability we are

T ₁	T ₂	T ₃
100 R(A)		
4=A-40 20 W(A)	60 A=A-20 W(A) 0	

T ₃	T ₂	T ₁
100 R(A)		
60 W(A) 20 W(A) 0		

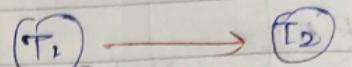
final result is equal
so both relats are equivalent

T ₁	T ₂
R(A)	
W(A)	R(A)
R(A) W(B)	W(A)
	R(B)

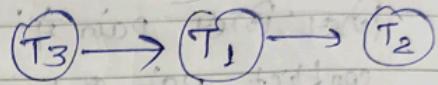


No cycle → conflict
serializable

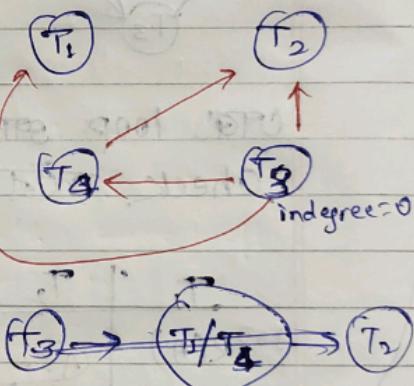
T_1	T_2	T_3 $R(N)$
$R(x)$		
$w(x)$		
	$R(y)$	
	$w(y)$	
		$w(x)$



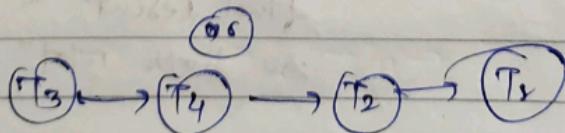
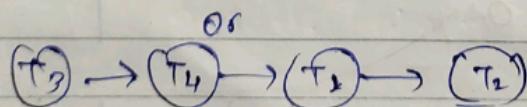
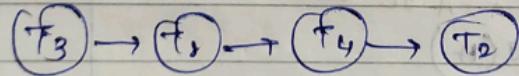
T_3
indegree = 0



T_1 $R(B)$	T_2	T_3	T_4
		$R(A)$	
		$w(A)$	
			$R(A)$
			$w(A)$
		$w(A)$	
			$R(c)$
			$w(c)$
$w(A)$			



means
Three possibility



Blind write : if there is no read happens prior to first write then it is said to be blind write.

T_1	T_2	T_3
	$R_2(x)$	
$R_1(x)$		$w_3(x)$
	$w_2(x)$	

$\rightarrow w_3(x)$ is blind write because there is no $R_3(x)$ before $w_3(x)$

$\rightarrow w_2(x)$ is also blind write

LC 10.3

VIEW SERIALIZABLE

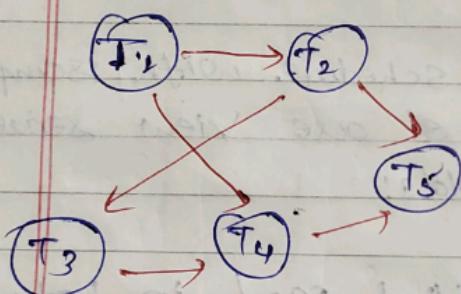
\rightarrow let S & S' be two schedules with same set of transaction. S & S' are view serializable if following cond' holds:

i) Initial read : initial read in both S & S' should be same.

ii) Final write : final write of each in both S & S' should be same

iii) Write Read pair : read perform by transact in both S & S' correspond to same write.

PA-10 Q-4	T_1	T_2	T_3	T_4	T_5
	$w(y)$	$r(y)$	$w(y)$	$r(y)$	$r(z)$
		$w(p)$	$w(y)$		$r(p)$
	$r(q)$			$w(x)$	
					$r(x)$ $w(y)$



→ acyclic

$T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_4 \rightarrow T_5$

possible serial conflict serialization

for view serializability

Check: initial read (i) final write (ii) write-read

(iii) write-read pairs:

• W-R pair $\rightarrow w(y) \rightarrow r(y) (T_1 \rightarrow T_2)$

$\rightarrow w(y) \rightarrow r(y) (T_3 \rightarrow T_4)$

$T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_4 \rightarrow (T_5)$
or

$T_3 \rightarrow T_4 \rightarrow T_1 \rightarrow T_2 \rightarrow T_5$

→ 2 view
serializability

#

RECOVERY:

- Recoverability is not possible if a transaction T_0 reads the updated item of T_1 another transaction and performs commit before T_1 .
- If T_0 first reads(A) then it must commit and

- # Cascading Rollback: A single transaction leads to failure of another transaction.

T_1	T_2	T_3	T_0
R(A)			(A)S
R(B)			(A)S (B)S
W(A)			(A)S
	R(A)		(A)S
	W(A)		(A)S
<u>abort</u>		R(A)	(A)C1

If T_1 fail then T_2 will also fail $\rightarrow T_3$ will also fail.

- # Cascadeless: To achieve cascadeless recovery a transaction T_1 must be committed before any other transaction reads the data written by T_1 .

Cascadeless \Rightarrow Cascading

→ every cascadeless is also recoverable.

Transact Control Language (TCL)

used with DML command

- COMMIT → to save changes
- ROLLBACK → to roll back the changes
- SAVE POINT → create points within group
- SET TRANSACTION → place a name on transaction

Q Check whether schedule is view serializable or not

T ₁	T ₂	T ₃
R(A)	R(B)	
R(A)	R(A)	
W(B)	R(A)	
	W(B)	
		W(B)

First read → T₂
last write → T₃

T₂ → T₁ → T₃

Here there is no write-read combination

tee
lockLOCK-BASED - PROTOCOL

→ lock is mechanism to control concurrent access to data item.

d) Exclusive (X) lock: both read + write

e) Shared (S) lock: only read

req. type →

		Share	exclusive
state of lock	Share	✓	X
	Exclusive	X	X

T₁T₂

Lock-s(A)

R(A)

Lock-s(B)

R(CB)

unlock(A)

unlock(CA)

Lock-X(A)

R(A)

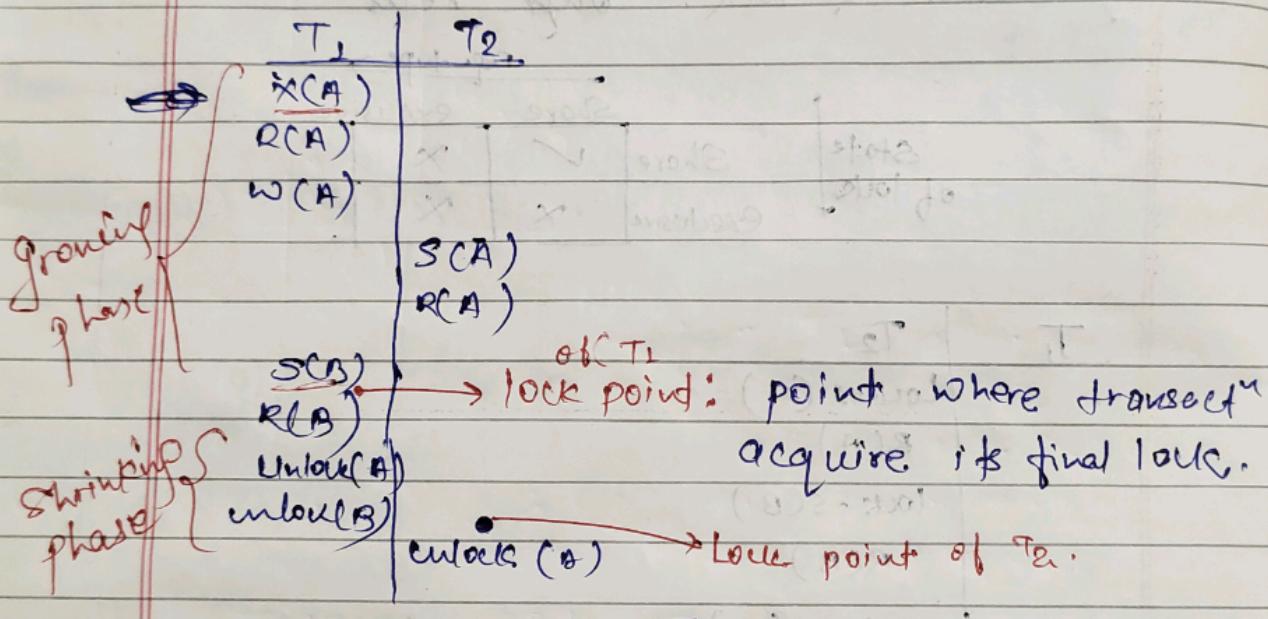
WCA)

unlock(A)

Two-Phase Locking protocol (2PL)

→ GROWING PHASE :- locks are acquired and no locks are released.

→ SHRINKING PHASE :- lock are released & no locks are acquired.



→ those Schedule which follow 2PL are always ~~conflict~~ serializable.

→ after shrinking phase growing phase is not possible

→ Growing phase is compulsory but shrinking phase is optional.

DEADLOCK:

→ occurs when two / more transaction are unable to proceed because each is waiting for the other to release resource.

T_1	T_2
lock-X(A)	lock-X(B)
RCA)	R(B)
WCA)	W(B)

T_1 want to update B
 T_2 want to update A.

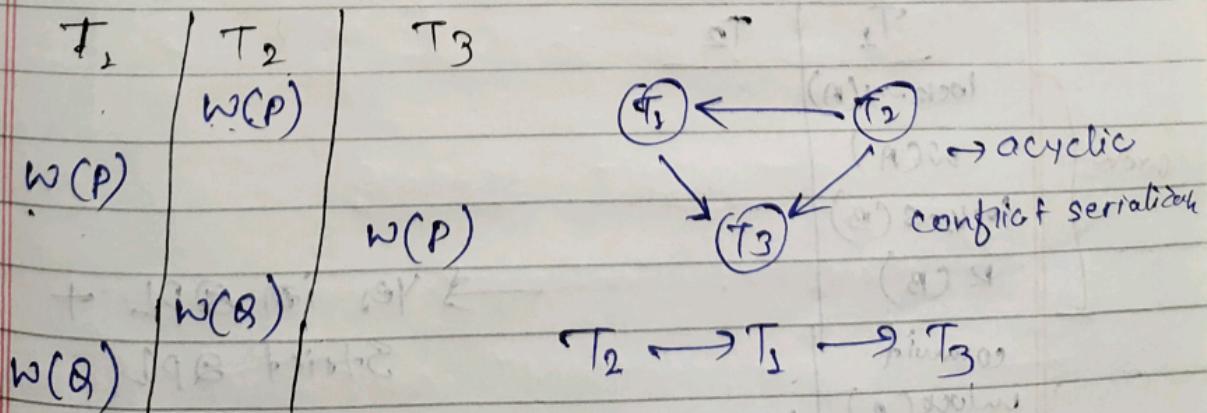
Here both T_1, T_2 are waiting for each other to give unlock
 this is deadlock.

→ if we do not unlock data item before requesting a lock on another data item, deadlock occurs.

→ deadlock is preferable than inconsistency state

Q. Schedule S : is given :

- vice versa
 - a) All conflict serializable schedule are 2D lock
 - b) give S is conflict serializable
 - c) given S is 2D lockable



T_1	T_2	T_3
$x(P)$		
$w(P)$		
$x(Q)$		
$w(Q)$		
$v(P)$		
$x(P)$		
$w(P)$		
$v(P)$		
$x(P)$		
$w(P)$		
$v(P)$		
$x(Q)$		
$w(Q)$		
$v(Q)$		
$x(Q)$		
$w(Q)$		
$v(Q)$		

(a) x-w-w (b) x-x-w
 (c) w (d) x
 (e) w-w (f) w-w

growing {
shrink }

growing {
shrink }

for T_1 , after growing
 phase one more growing
 phase came. $\times \times$

for T_2 , after ~~growing~~^{shrinking}
Phase one more growing
phase came. X X

not follow 2PL

STRICT 2PL

→ follow QPL and all exclusive lock.
Should unlock after commit operation.

	T_1	T_2
grow	lock-X(A)	
	WCA)	>
	lock-S(B)	
	R(B)	
	commit	
	unlock(A)	
shrink		
	lock-X(A)	
	WCA)	grow
	commit	
	unlock(A)	shrink
	unlock(B)	

Rigorous & PL

↳ follow & PL

↳ all lock (exclusive/shared) should outcome after commit.

lock manager:

↳ process to which a transaction sends lock and unlock request.

↳ reply a lock request by sending message asking transaction to roll back, (if dead lock happens)

Deadlock, Timestamp:

→ each "transact" that enters the system is assigned a unique time stamp, typically based on order they arrive.

① System clock

② logical counter (1, 2, 3, 4, ...)

Wait-Die Scheme

↳ non-preemptive technique

→ older "transact" may wait for younger one to release data (Older means small timestamp)

Ex suppose T_5, T_{10}, T_{15} have time stamp 5, 10, 15 resp
 $T_5 \rightarrow$ older one (small timestamp)

then,

→ if T_5 request a data held by T_{10} then
 T_5 will wait.

→ if T_{15} request data held by T_{10} then
 T_{15} will be killed

Conc

Wound-Wait Scheme

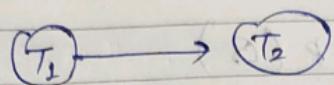
↳ preemptive technique

→ younger "transact" may wait for older one in last ex:

→ if T_5 request a data held by T_{10} then T_{10} will be wounded.

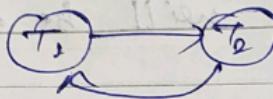
→ if T_{15} request data held by T_{10} , then
 T_{15} will wait.

Deadlock detection



means T_2 is waiting for T_1 to release data.

→ System is in deadlock if cycle occurs.



both are waiting for each other to release data.

For deadlock Recovery:

Some transaction have to rolled back to break deadlock.

this is done through:

Time stamp-Based Protocol

→ W-timestamp(α) → largest time taken by any transaction to execute $W(\alpha)$ successfully.

→ R-timestamp(α) → largest time taken to execute $R(\alpha)$.

Case T: If T_1 issues a read operatⁿ $R(\alpha)$ then:
 $TS(T_1) \leq W\text{-timestamp}(\alpha)$
then T_1 will be rejected.

$TS(T_1) \geq W\text{-timestamp}(\alpha)$

then T_1 will execute

~~core-1~~ T_1 issues $w(\varnothing)$ then:

\rightarrow If $TS(T_2) < R\text{-timestamp}(\varnothing)$

$\hookrightarrow T_2$ will be rejected

\rightarrow If $TS(T_2) < w\text{-timestamp}(\varnothing)$

$\hookrightarrow T_2$ will be rejected

\rightarrow If $TS(T_2) \geq R/w\text{-timestamp}(\varnothing)$

$\hookrightarrow T_2$ will be executed

~~Q8~~ S: $R_1(A), R_2(B), W_2(A), R_1(A)$

TS for T_1 & T_2 are 1, 2 respec.

a) T_1 need to roll back $TS(T_1) \leq 1$

b) T_2 need to roll back $TS(T_2) \geq 2$

$$R_2(A) = 0$$

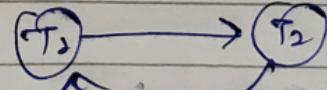
T_1	T_2
$R(A)$	
$R(B)$	

$W_2(A) = R(A)$

$1 < 2$

roll back

	A	B
W-TS	2	
R-TS	1	1



deadlock occurs
one transaction need to
roll back: we have
to find which one
to roll back.